

3+1D GNLSE Solver Documentation

This code is currently suitable for forward simulation of pulses and CW propagation. **While the code is autodifferentiable (written in JAX), and theoretically can be plugged into optimizers, this current distribution of the code contains no additional infrastructure for doing so.** This code solves a version of the Unidirectional Pulse Propagation Equation (UPPE), a (3+1)D Generalized Nonlinear Schrödinger Equation (GNLSE):

$$\begin{aligned} \frac{\partial A(x, y, z, \omega)}{\partial z} = & \frac{1}{2\beta_{eff}(\omega)} \nabla_{\perp}^2 A + \hat{D}_{\omega}(\omega) A + i \frac{\beta_{eff}(\omega)}{2} \left(\left(\frac{n(x, y, \omega)}{n_{eff}(\omega)} \right)^2 - 1 \right) A \\ & + \mathcal{F}_t \left\{ i \frac{n_2 \omega_0}{c} \left(1 + \frac{i}{\omega_0} \frac{\partial}{\partial t} \right) \{ (1 - f_R) |A|^2 A + f_R A [h * |A|^2] \} \right\} \end{aligned}$$

The RHS terms describe, from left to right, *diffraction*, *dispersion*, *waveguiding*, and finally nonlinear effects including *Kerr nonlinearity* and *Raman scattering* (tuned between with the *Raman fraction* $f_R \in [0, 1]$), and self-steepening. Other mechanisms implemented in code (but not indicated in the above equation) include *gain*, a *saturation intensity*, and perfectly matched layers (PMLs) that prevent undesired back-reflections at simulation boundaries.

Waveguiding, diffraction, dispersion, Kerr nonlinearity, and PMLs have been tested and appear to operate reliably; other effects (e.g., self-steepening, Raman scattering, gain, etc.) have not been thoroughly tested--use at your own risk (or else report anything fishy to me).

Quick Guide to Turning off Effects:

Setting arguments/flags to exactly these values not only ensures they have zero magnitude, but also signals to the code to completely skip those calculations (and speed up sim) where possible.

- Raman Scattering off: turn $f_R = 0.0$.

To construct and run a simulation requires specification of the following:

- Source: An initial field of the form $A(x, y, t)$, launched at the $z = 0$ plane of the propagation medium.
- Medium: defined by the refractive index profile $n(x, y, \omega)$, nonlinear coefficient n_2 , and material dispersion parameters (β_1, β_2) . (Note: β_0 is determined by the source wavelength and the refractive index.) Additional effects might include Raman response (parameterized by t_1 and t_2 , turned "off" by setting), self-steepening effects

gnlse_solver.py: Runs the simulation. The only user-facing function should be

- `GNLSE3D_propagate(args, A0, *, event_fn=None, event_payload=None, stop_on_event=True, event_check_every: int=1) → dict(field, dt, dx, seconds, **meta)`
 - `args` : a dictionary containing all physical information about the simulation, including:
 - `"Lx", "Ly", "Lz", "Lt"` : Physical dimensions $[m]$, $[s]$.
 - `"Nx", "Ny", "Nt"` : Grid sizes, ints.
 - `"deltaZ", "deltaZ_NL"` : Linear and nonlinear step-sizes $[m]$.
 - `"save_at"` : Array of z-locations to save $A(x, y, t)$.
 - `"lambda0"` : Central wavelength of pulse or CW source in free space $[m^{-1}]$.
 - `"n2"` : Nonlinear coefficient
 - `"beta0", "beta1", "beta2"` : Dispersion coefficients.
 - `"gain_coeff", "gain_fwhm"` : Gain parameters (set coeff to 0.0 to turn off).
 - `"t1", "t2"` :
 - `"n_xyomega"` : Refractive index profile, specified for each frequency.
 - `"pml_thickness"` : depth of PMLs, in $[m]$.
 - `"pml_wmax"` : maximum damping exponent of the PML, in $[W/m]$.
 - `"m_nl_substeps"` : Number of substeps used in the nonlinear step of integration.
 - `"A0"` : The initial field at $z = 0$, $A(x, y, t)$. Has dimensions $N_x \times N_y \times N_z$.

gnlse_medium.py: Contains functions for creating the refractive index profile $n(x, y, \omega)$.

Refractive index profiles:

- `make_space(Lx, Nx, Ly, Ny) → X, Y`
 - `Lx, Ly` : spatial x/y dimensions of sim.
 - `Nx, Ny` : spatial x/y grid sizes.
 - `X, Y` : arrays of physical distances along axes.
 - `dim(X) = Nx, dim(Y) = Ny`.
 - Use the outputs of `make_space` as inputs `X, Y` of the below functions.
- `make_polynomial_n(X, Y, n_core, n_clad, r_core, alpha = 2) → n(x, y):`
 - Returns $n(x, y) = n_{\text{core}} \sqrt{1 - 2\delta \left[\frac{x^2 + y^2}{r_{\text{core}}^2} \right]^\alpha}$ inside core radius, and n_{clad} without.
 - $n(x, y)$ is an $N_x \times N_y$ dimensional array.
 - $\delta = \frac{n_{\text{core}} - n_{\text{clad}}}{n_{\text{core}}}$
 - E.g., use $\alpha = 2$ for GRIN fiber.

- `make_supergauss_n(X, Y, n_core, n_clad, r_core, m=20) → $n(x, y)$`
 - $n(x, y) = 1 + \dots$
- `make_bulk_n(X, Y, n) → $n(x, y)$` : Unwritten, will wrap `n * jnp.ones((len(X), len(Y)))`
- `bend_n(n_xy, Lz, a) → $n(x, y, z)$` : Unwritten
- `gaussian_disorder_n`: Unwritten.

gnlse_source.py: Contains functions for building transverse and temporal field profiles.

Generating Modes:

- f
- Accessing Modes:

Generating Temporal Profiles:

gnlse_visualizations.py: Contains functions for plotting simulation outputs. These are heavily GPT'ed since they're pretty trivial matplotlib tasks. Presumably the user will DIY most of their own plotting.

gnlse_events.py: Contains functions for implementing event handling in simulations (e.g., halting simulation at onset of self-focusing collapse).

Ziyu recommended a Basler ace camera (this one seems suitable:

<https://www.baslerweb.com/en/shop/aca3088-57um/>) and a pco.edge 4.2 camera, if we want higher resolution or lower noise: <https://www.excelitas.com/product/pcoedge-42-usb-scmos-camera>. **I think the Basler camera seems adequate for what we're doing, but we can discuss.** The Basler ace A3088-57um has:

- 7.41 mm x 4.95 mm sensor area.
- 3088 x 2064 (6 MP) resolution.
- 2.4 micron x 2.4 micron pixel size.
- 59 fps framerate.

Ziyu said he's using a 600 x 800 pixel Basler, which I think has higher framerate, but (discussed below), I'm not sure (a) we need a framerate higher than what we expect the SLM's to be and (b) I think we have a premium on spatial resolution since we're looking at fine features of the transverse field.

My intuition for specs is probably a little skewed, since my experience is with applications requiring single-photon resolution. That being said, my assumptions for our use-case are as

follows:

- Pixel pitch: two options for reference are either the scale of the SLM pixels (which, I think, are ~ 10 s of microns), or (more demanding) the order of 2-4 times the grid used in simulation (since the pixel-scale features in simulation are somewhat unreliable anyway). So I think higher spatial resolution is actually something to prioritize here, especially as some of our optimization (or Novelty Search) targets may involve "shrinking" feature sizes.
- Quantum efficiency of visible spectrum cameras at 1064 nm is nonzero, but low--but I can't imagine we are at a shortage of photons here.
- I do not have a good sense of how noise works for these high power cases; I don't think we're doing anything so subtle that I can imagine it being a big consideration here, but the pco.edge 4.2 is an option if it's important.
- Assuming our worst-case-scenario output profile dimensions are simply those of ~ 10 cm linear propagation (diffraction) of a 3 mm diameter, 1064 nm beam, the broadening is effectively negligible. Any sensor of dimensions ≥ 3 mm by 3mm should be fine--probably at least 5mm by 5mm (without sacrificing resolution) would make life easier.
- The nonlinear propagation distance required will depend on the wedge/block material. Dongjin and I are going to run some bulk media propagation demos for his presentation, so we can double check this 10 cm estimate quickly. However, I strongly doubt it will change this beam-broadening estimate drastically.
- Unless I am misunderstanding, I don't think sub-pulse duration-resolving framerates are reasonable or expected here. I imagine our intention is to send one-shape-of-pulse-per-frame, limited by the slower of either the framerate or SLM speed. SLM framerate (from Thorlabs, as reference) is ~ 60 Hz. So we can probably get away with a low framerate on the camera. It seems like we'd want $r_{slm} \geq r_{cam}$ anyway, for things to be easily synchronized.